

◆ WEB3 SMART CONTRACT SECURITY REVIEW

REWARD ACCOUNTING DRIFT

Solidity · Staking · Reward
accounting

Manual review + working proof-of-concept +
severity-rated report

● High 1



GILLES MUSY

Smart contract security researcher

Confirmed findings · Code4rena GiMu84 · Cantina GiLMu

1 Disclaimer

This document is a demonstration security review performed on RewardPool, a deliberately vulnerable sample authored by Gilles Musy for portfolio purposes. It is not a client engagement: the code was never deployed, holds no user funds, and the vulnerabilities it contains were intentionally introduced to illustrate review methodology, severity classification and proof-of-concept construction. The findings, proofs of concept and remediation guidance below are genuine analysis of the code as written.

A security review is a time-, resource- and expertise-bound effort to surface as many issues as possible; it cannot prove the absence of vulnerabilities. Subsequent reviews, a bug-bounty program and runtime monitoring are recommended for any production deployment.

2 Risk classification

Severity is derived from the impact of an issue and the likelihood of its occurrence.

IMPACT ▼ / LIKELIHOOD ►	HIGH	MEDIUM	LOW
HIGH	Critical	High	Medium
MEDIUM	High	Medium	Low
LOW	Medium	Low	Low

IMPACT High: significant loss of assets or harm to many users. Medium: moderate loss or harm. Low: minor loss or harm.

LIKELIHOOD High: practical under realistic on-chain conditions at low cost. Medium: conditionally incentivised but plausible. Low: needs unlikely assumptions or significant attacker stake.

3 About the target

RewardPool, a MasterChef-style single staking pool.

TYPE Demonstration review on intentionally vulnerable code

METHOD Foundry proof-of-concept, vulnerable branch (master) and remediated branch (fixed)

4 Findings

HIGH H-01, Stale accumulator on deposit (High)

A depositor who joins a staking pool right before its reward accumulator is refreshed is credited for rewards that accrued while they held no stake. In the proof-of-concept Alice is the sole staker for 100 seconds and earns all 100 reward tokens of that window; the attacker deposits at the end, harvests immediately, and takes 50 of them. The 50 tokens come straight out of Alice's earned rewards. The attack needs no privilege and no interaction from the victim.

ROOT CAUSE

The pool uses the standard MasterChef accounting: `accRewardPerShare` accumulates rewards per staked unit, and each user's `rewardDebt` records the accumulator value at their last interaction, so `pending = amount * accRewardPerShare - rewardDebt`. The invariant is that `accRewardPerShare` must be brought current before any stake changes. `deposit` breaks it:

```
function deposit(uint256 amount) external {
    UserInfo storage u = users[msg.sender];
    if (u.amount > 0) { ... }
    stakeToken.transferFrom(msg.sender, address(this), amount);
    u.amount += amount;
    totalStaked += amount;
    u.rewardDebt = (u.amount * accRewardPerShare) / ACC; // stale accRewardPerShare
}
```

No `updatePool()` runs first. When rewards have accrued since the last refresh, `accRewardPerShare` is still the old, lower value. The new depositor's `rewardDebt` is snapshotted against that stale value, so the next refresh credits them for the gap.

ATTACK

1. Alice deposits 100 at time `t0`. She is the sole staker. `accRewardPerShare` is 0.
2. 100 seconds pass. 100 reward tokens accrue, all owed to Alice. The accumulator has not been refreshed, so it still reads 0 and `lastRewardTime` is still `t0`.
3. Attacker deposits 100 at `t0 + 100`. No `updatePool` runs, so the attacker's `rewardDebt` is set against the stale accumulator value of 0. `totalStaked` becomes 200.
4. Attacker calls `harvest`, which refreshes the accumulator: the 100 accrued tokens are now spread over `totalStaked = 200`, giving `accRewardPerShare = 0.5` per unit. The attacker's pending is `100 * 0.5 - 0 = 50`. The attacker harvests 50.
5. Alice harvests `100 * 0.5 - 0 = 50` instead of the 100 she earned.

PROOF OF CONCEPT

`test/RewardPool.poc.t.sol`, function `test_accountingDrift_stealsAccruedRewards`, run on master:

```
rewards accrued while only Alice staked (wei): 100000000000000000000
attacker harvested (wei): 50000000000000000000
alice harvested (wei): 50000000000000000000
```

RECOMMENDATION

Call `updatePool()` at the start of `deposit`, before the stake or `rewardDebt` is touched:

```
function deposit(uint256 amount) external {
    updatePool();
    UserInfo storage u = users[msg.sender];
    ...
    u.rewardDebt = (u.amount * accRewardPerShare) / ACC; // now current
}
```

Refreshing the accumulator first snapshots `rewardDebt` against the up-to-date `accRewardPerShare`, so a new depositor is credited only from the moment they join. Every stake-changing entry point (`deposit`, `withdraw`, `harvest`) must refresh the accumulator before mutating state. On the `fixed` branch the same proof-of-concept shows the attacker harvesting 0 and Alice keeping the full 100.

SEVERITY

High. The impact is direct theft of accrued rewards from honest stakers. The actor needs no privilege, and the only precondition, depositing while rewards are pending accrual, is the normal state of an active pool.

5 Informational, Gas & Non-Critical

NON-CRITICAL N-01, No events on state changes (Non-Critical)

`deposit`, `harvest` and `updatePool` (L50, L62, L38) mutate stakes and pay rewards without emitting any event, leaving pool activity untraceable off-chain. Add `Deposit`, `Harvest` and `PoolUpdated` events.

INFORMATIONAL I-01, Unchecked ERC20 return values (Informational)

`rewardToken.transfer` (L54, L66) and `stakeToken.transferFrom` (L56) ignore the boolean return value. A non-reverting, non-conforming token would fail silently and desynchronise accounting. Use OpenZeppelin `SafeERC20`.

INFORMATIONAL I-02, Reward-paying paths are not reentrancy-guarded (Informational)

`deposit` and `harvest` (L50, L62) make external token calls while mutating user accounting and carry no `nonReentrant` guard. Independently of the main accumulator finding, add a reentrancy guard as defense in depth.

Scope and disclaimer

`RewardPool` is intentionally vulnerable code written to demonstrate audit methodology end to end. It is not production code and must never be deployed. The finding above is a real vulnerability in this demo contract, reproduced with an executable proof-of-concept, not an invented severity.